

# Retrieval Augmented Generation RAG Improvements



Building Competence. Crossing Borders.


Jasmin Heierli

jasmin.heierli@zhaw.ch, 20 March 2024

# Cosine Similarity

Let's close the circle and go back to our original Problem:

We can now calculate the cosine similarity between two vectors



$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i \cdot B_i}{\sqrt{\sum_{i=1}^n A_i^2} \cdot \sqrt{\sum_{i=1}^n B_i^2}}$$

$$\mathbf{A} \cdot \mathbf{B} = 50 \times 48 + 10 \times 12 = 2400 + 120 = 2520$$

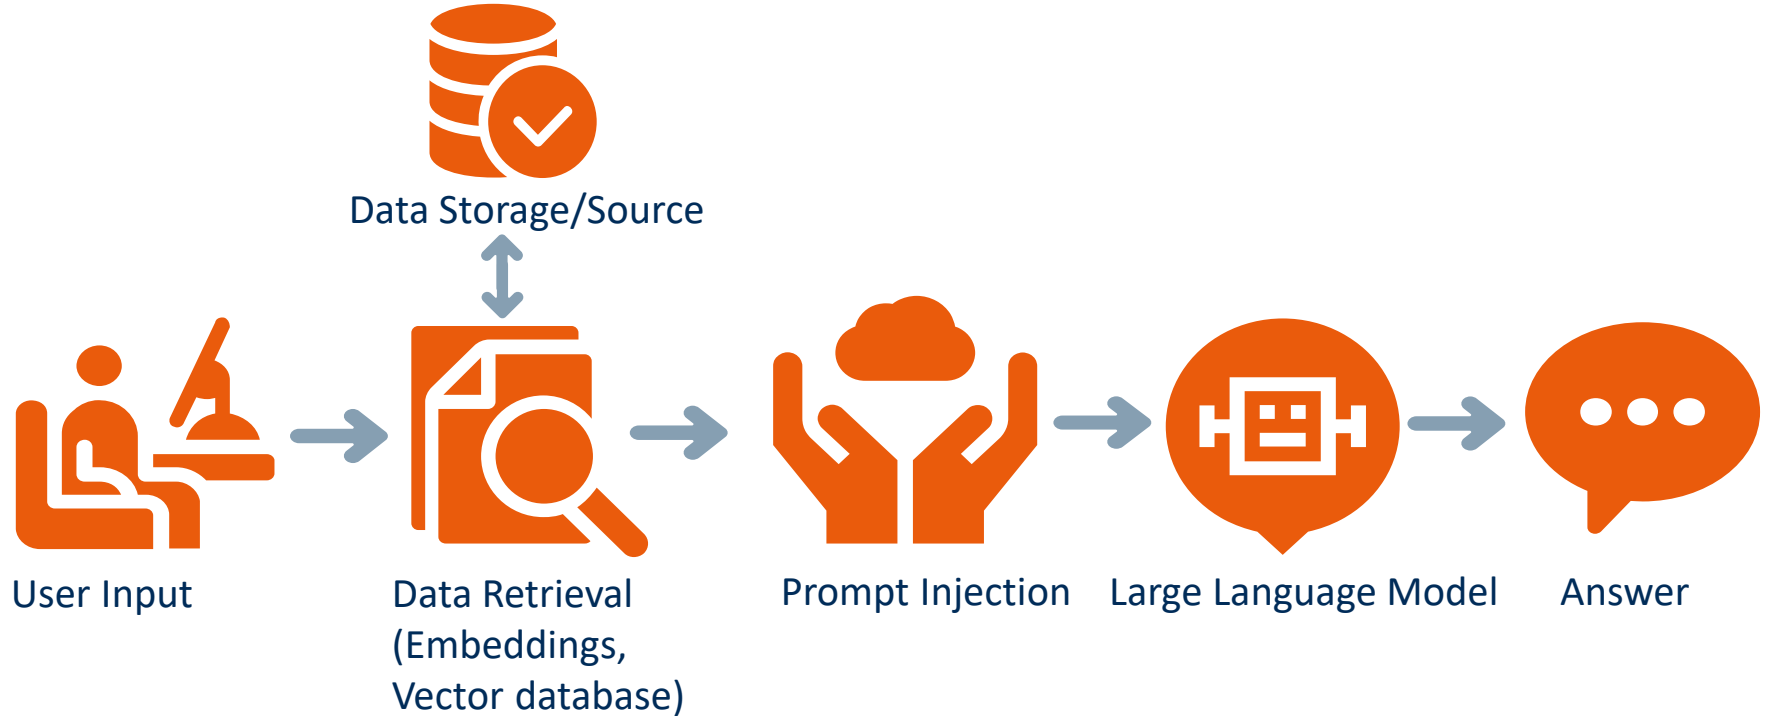
$$\|\mathbf{B}\| = \sqrt{48^2 + 12^2} = \sqrt{2304 + 144} = \sqrt{2448} \approx 49.5$$

$$\|\mathbf{A}\| = \sqrt{50^2 + 10^2} = \sqrt{2500 + 100} = \sqrt{2600} \approx 51.0$$

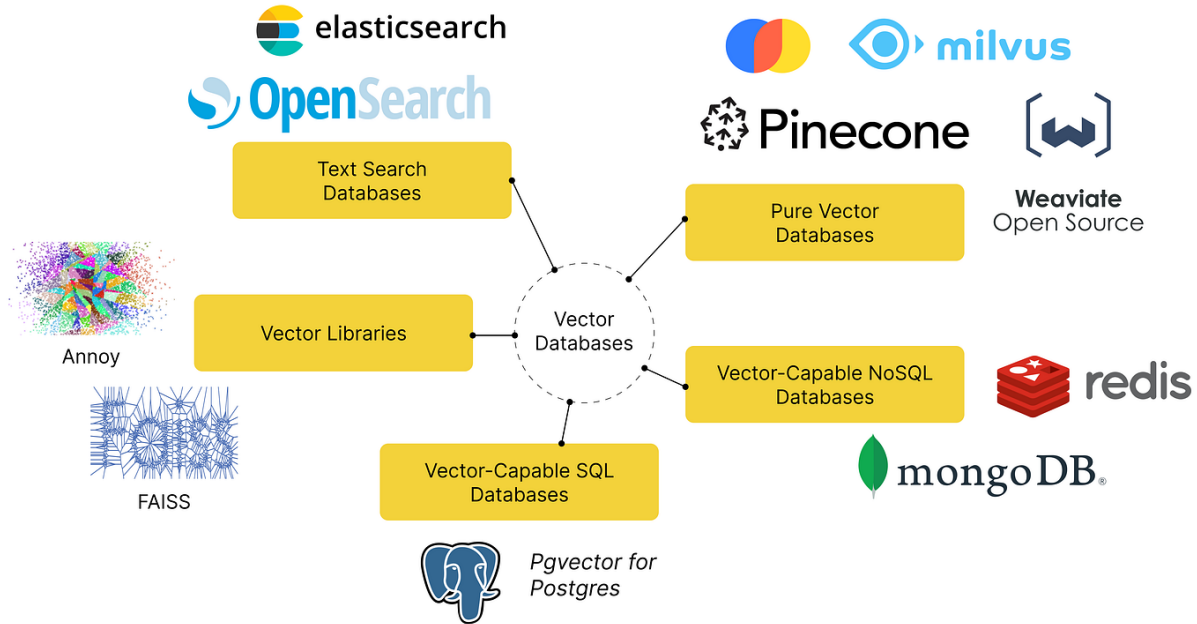
$$\text{cosine similarity} = \frac{2520}{51.0 \times 49.5} \approx 0.998$$

$$\cos(2.73^\circ) \approx 0.998$$

# RAG Architecture



# Vector Stores



<https://towardsdatascience.com/all-you-need-to-know-about-vector-databases-and-how-to-use-them-to-augment-your-llm-apps-596f39adfedb>

# Vector Store vs. „Traditional“ Database

Vector Store	Traditional Database
Indexed unstructured data	Structured data in tables
Results based on similarity	Results based on exact keyword match
Store text/images and their embeddings	Store in database scheme pre-defined attributes
Index groups similar documents	Table groups related attributes

<https://towardsdatascience.com/explaining-vector-databases-in-3-levels-of-difficulty-fc392e48ab78>

# Vector Stores Indexing

- **Problem statement:** Calculating the similarity between your vector and every database entry is expensive and time-consuming
- Vectors are **indexed**: they are organized to enable an efficient data retrieval.
- Solution: **approximate nearest neighbor (ANN)** approach
  - Pre-calculate distances between vector embeddings and organise and store similar vectors close to each other
  - Vectors that are similar to each other are stored close to each other in clusters or a graph

We trade in accuracy for speed.

# ANN Algorithms

## Overview Only:

- Clustering-based index
  - FAISS
- Proximity graph-based index
  - HNSW
- Tree-based index
  - ANNOY
- Hash-based index
  - LSH
- Compression-based index
  - PQ

<https://weaviate.io/blog/why-is-vector-search-so-fast>

# Problem: Our retrieved documents do not seem relevant enough

Given cosine similarity for the retriever

Give our database about air travel

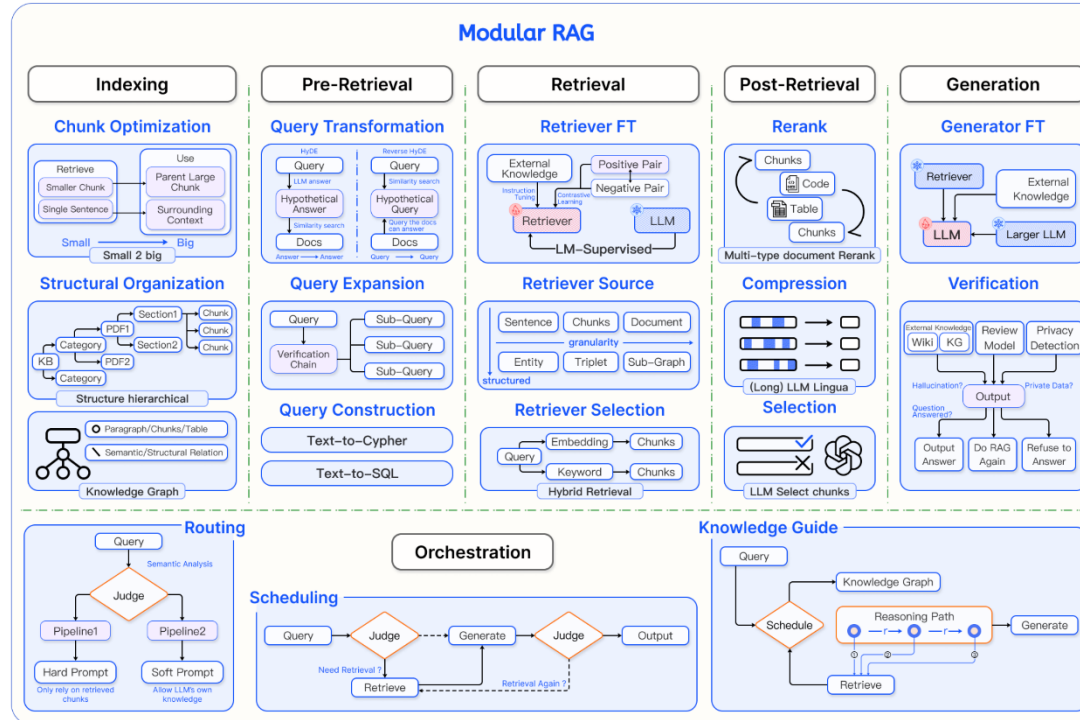
Given our original query

**What were the first words spoken after the landing on the moon?**

**The first landing on the moon was on 21 July 1969 according to a spokesperson...**



# Potential Improvements



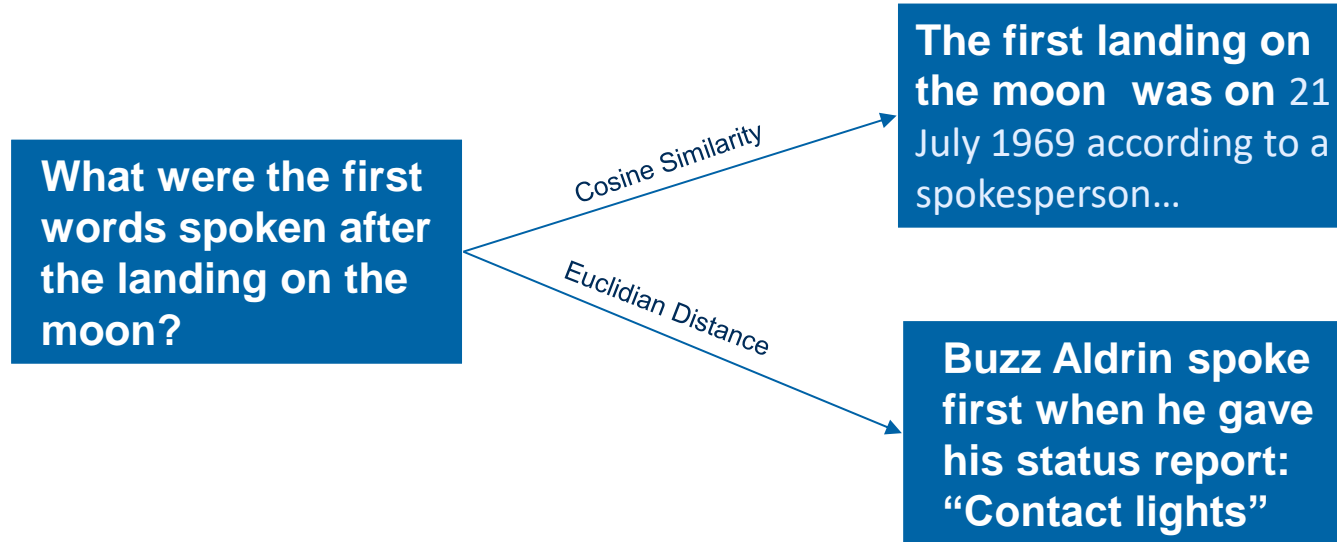
<https://towardsdatascience.com/implementing-modular-rag-with-haystack-and-hypster-d2f0ecc88b8f/>

# Embedding Adaptors

- Goal: transform query embeddings that they align better with the target task's feature space
- **Adapter**: an alternative to fine-tune an entire pre-trained model, which is implemented as small feed-forward neural networks between layers of pre-trained models
- Advantage: can sustainably improve results
- Disadvantage: requires user feedback/training data for relevant documents



# Choice of Similarity Metric



<https://medium.com/@stepkurniawan/comparing-similarity-searches-distance-metrics-in-vector-stores-rag-model-f0b3f7532d6f>

# Hybrid Search

**What were the first words spoken after the landing on the moon?**

Vector Similarity


**The first landing on the moon was on 21 July 1969 according to a spokesperson...**

**What were the first words spoken after the landing on the moon?**

Key Words: First words spoken

**Buzz Aldrin spoke first when he gave his status report: "Contact lights"**

# Distance Thresholding



- Classic idea behind top k retrieval
- Decide on a max. acceptable distance between query and chunk vector
- Threshold can be used exclusively or combined with top k retrieval
- Idea: Find maximum possible proximity with predetermined question and answer pairs and maximum negative question and answer pairs with unrelated examples to find a good threshold

# Structure of Base Prompt

- Think about where to inject the retrieved information
- Think about the structure of the injected information
  - Consider using simplified XML tags
  - Information at the beginning and at the end of a prompt is more likely to be considered


**Answer the following question: {question}**

Base your answer solely on the information given:

<information> {information} </information>

# Query Transformation: Query Expansion

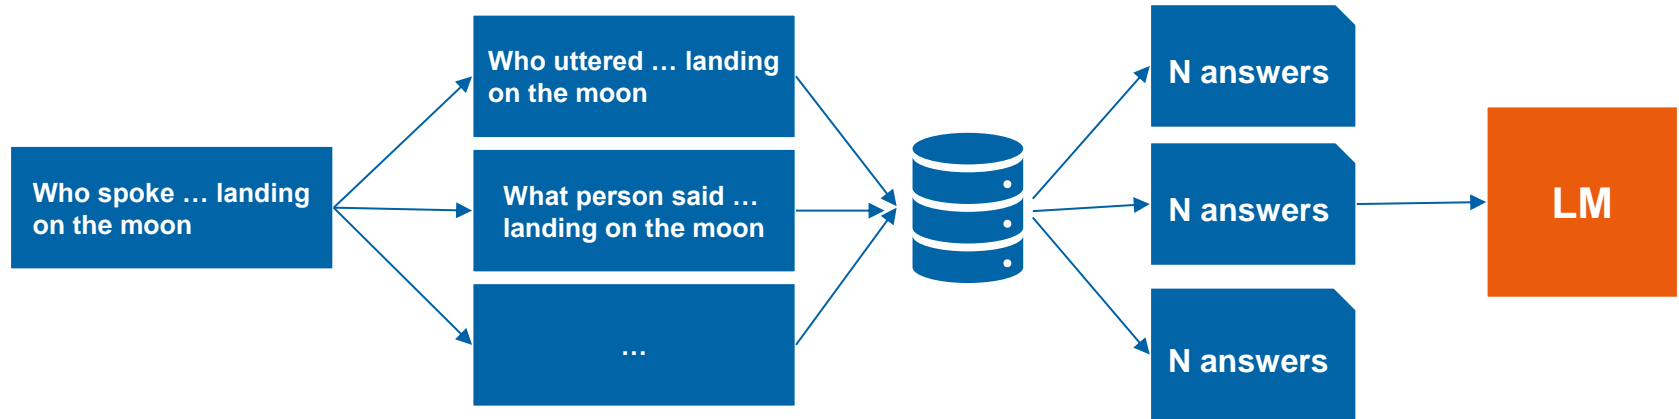
- Expand the query with a hypothetical answer for the question
- Does not really matter whether the answer is correct or not
- Based on the hypothetical answer the result is expected to be better
- The original query may or may not be sent too



<https://towardsdatascience.com/3-advanced-document-retrieval-techniques-to-improve-rag-systems-0703a2375e1c>

# Query Expansion: Query Augmentation / Rewrite

- Ask an LM to rewrite the given query → humans sometimes phrase their questions badly
- Ask the LM to generate queries that are related to the query, ask like an expert would ask, a 5 year old...
- Also works given just a keyword, depending on the prompt for the LM

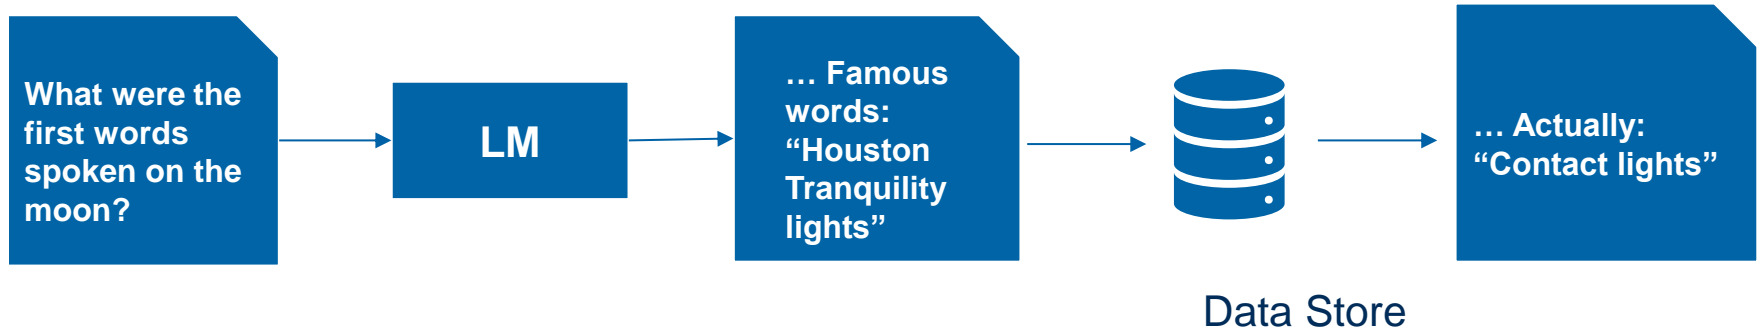


<https://towardsdatascience.com/9-effective-techniques-to-boost-retrieval-augmented-generation-rag-systems-210ace375049>



# Query Transformation: HyDe

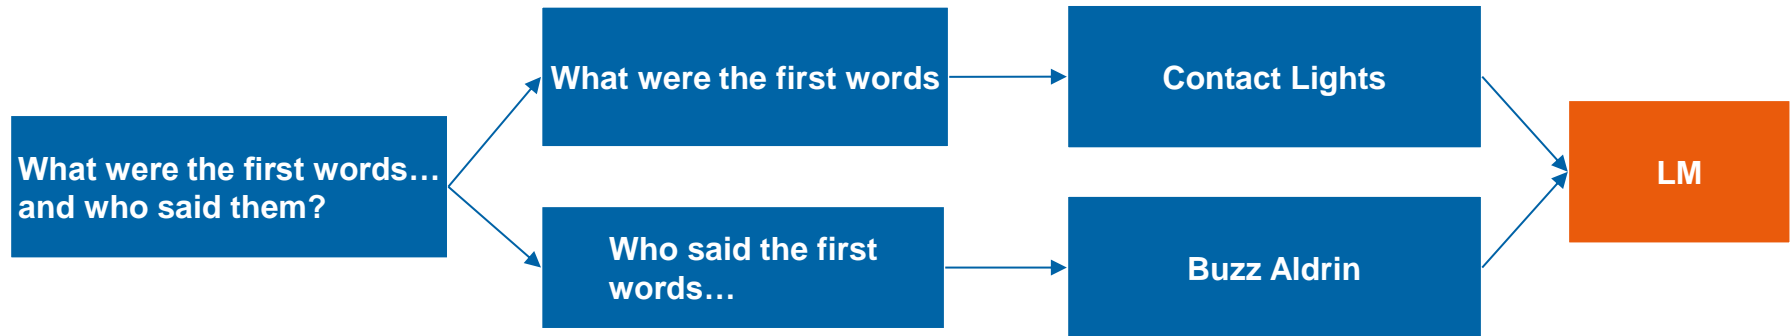
- **HyDe**: Hypthetical Document Embeddings
- We prompt the LM to generate/hallucinate a paragraph that would answer our question
- We send the hypothetical document to our retriever



<https://towardsdatascience.com/3-advanced-document-retrieval-techniques-to-improve-rag-systems-0703a2375e1c>

# Query Transformations: Subquestions/Query Planning

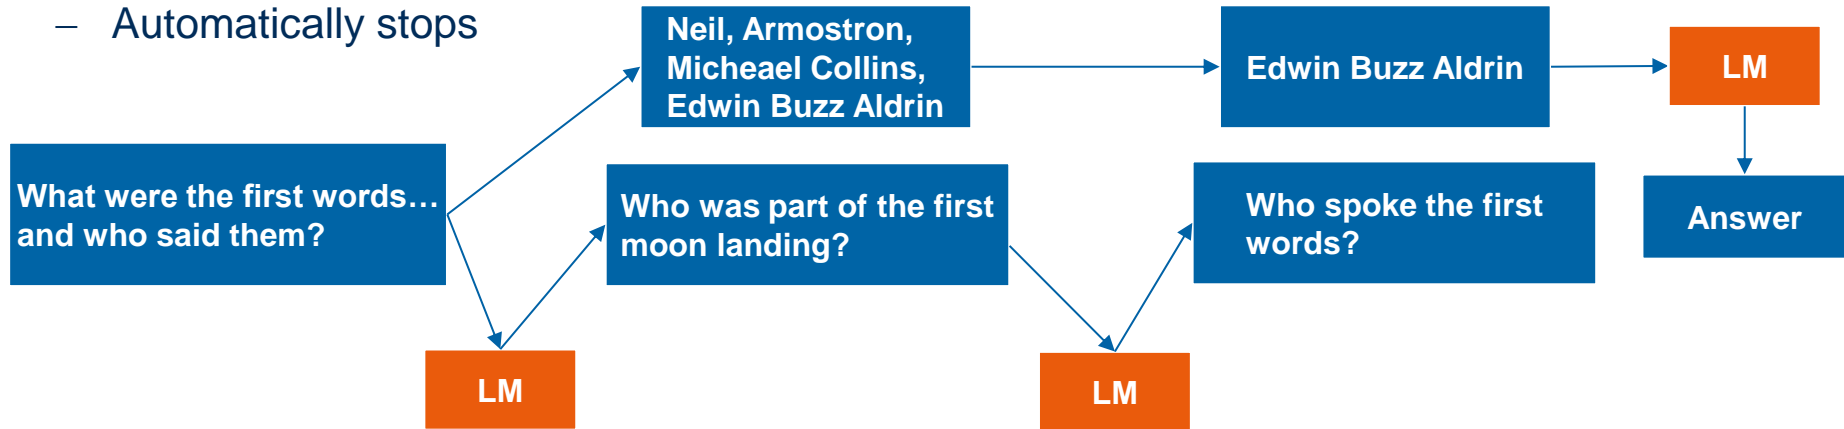
- Divide and Conquer approach to answer complex questions
- Analyze the original question and break them down into smaller subquestions
- Answers to subquestions are concatenated
- Concatenated answers will be used to generate the final question



<https://towardsdatascience.com/advanced-query-transformations-to-improve-rag-11adca9b19d1>

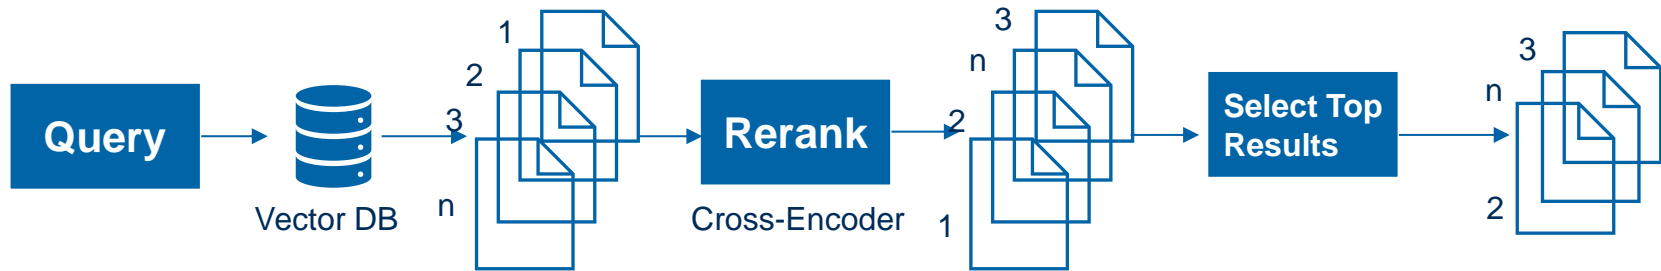
# Query Transformation: Multi-Step Query Transformation

- Based on the self-ask / reasoning method
- LM is asked whether a follow-up question is required to answer the original query
- After each answer, the model will be asked whether it requires more subquestions
- Automatically stops



# Cross-encoder Reranking

- “Second opinion” on the relevancy of results
- Retrieve more documents than required
- Cross-encoder compares each result with the query
- Sorts the result in decreasing order



<https://towardsdatascience.com/3-advanced-document-retrieval-techniques-to-improve-rag-systems-0703a2375e1c>

# End-to-end Evaluation of RAGs

- Generate a set of questions (min 30-50) for a subset of documents with an LM
  - Ideally, generate more questions than you need and remove duplicates or add more relevant examples
- Let your RAG application answer all of these questions and store the answer alongside the original question
- Let a high-quality LM (GPT-4 for example) evaluate whether the answer is relevant or not
- Manually double-check the judgments

```
"A sample from the documents is below.\n"\n"-----\n"\n"{context_str}\n"\n"-----\n"\n"Using the documentation sample,\n"{query_str} generate more questions similar to the\nexample"
```

<https://betterprogramming.pub/exploring-end-to-end-evaluation-of-rag-pipelines-e4c03221429>

# Thank you.

