

Machine Learning II

Generative AI



Dr. Elena Gavagnin

elena.gavagnin@zhaw.ch

FS2025

Machine Learning II – Images

Plan:

1. Artificial Neural Networks – Introduction, implementation with TF/Keras and Training
2. "Classic" Deep Learning & Images – Convolutional Neural Networks and Transfer Learning
3. Multimodal and generative AI

«Multimodal and generative AI (W10 and W11):

1. Extracting Infos from Images: Object Detection
2. Multimodal AI: Connecting Anything and Images (CLIP, OWL-ViT)
3. Foundational Models in CV
3. Generative AI: from Autoencoder to Diffusion Models (and past that)


Generative AI: from Autoencoders to Diffusion Models

Efficient Data Representations

Which of the following number sequences do you find the easiest to memorize?

- 40, 27, 25, 36, 81, 57, 10, 73, 19, 68


- 50, 48, 46, 44, 42, 40, 38, 36, 34, 32, 30, 28, 26, 24, 22, 20, 18, 16, 14



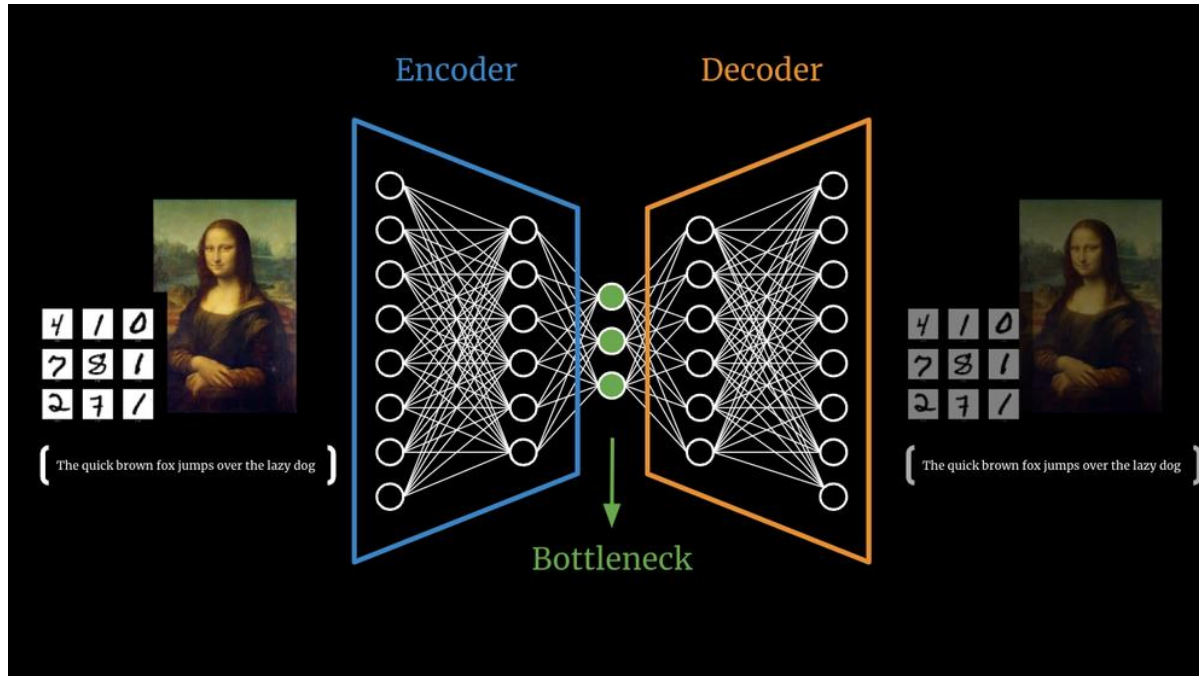
Autoencoders

Autoencoders = artificial neural networks capable of learning dense representations of the input data, called *latent representations* or *codings*, without any supervision (i.e., the training set is unlabeled).

much lower dimensionality! Useful for dimensionality reduction, feature detectors (unsupervised pretraining of NN) or as generative models:




Autoencoders



number of neurons in the output layer must be equal to the number of inputs.

Autoencoder



```
stacked_encoder = tf.keras.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(30, activation="relu"),
])
```



```
stacked_decoder = tf.keras.Sequential([
    tf.keras.layers.Dense(100, activation="relu"),
    tf.keras.layers.Dense(28 * 28),
    tf.keras.layers.Reshape([28, 28])
])
```

```
stacked_ae = tf.keras.Sequential([stacked_encoder,
                                   stacked_decoder])
```

```
stacked_ae.compile(loss="mse", optimizer="nadam")
```

```
history = stacked_ae.fit(X_train, X_train, epochs=20,
                        validation_data=(X_valid, X_valid))
```


Autoencoders



- Learned features can be used in pre-training
- Can combine with Convolutions
- Force Learning: Denoising Autoencoders

VAE: Variational Autoencoders


<https://arxiv.org/abs/1312.6114>

Specialty:

- *probabilistic* models
-> outputs are partly determined by chance, even after training
- *generative*:
-> generate new instances that look like they were sampled from the training set

How to generate new images?

Sample random encodings from a Gaussian distribution and decode them!



GANs: Generative Adversarial Networks

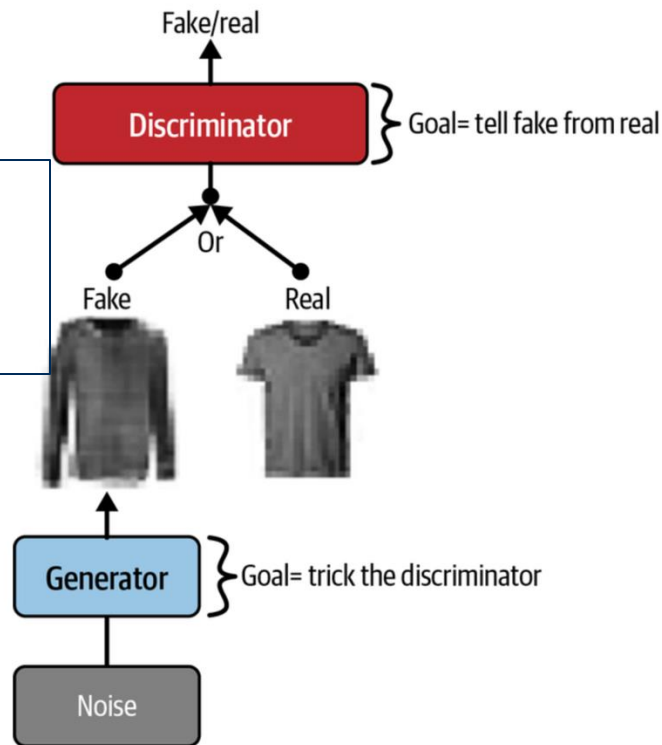
<https://arxiv.org/abs/1406.2661>

2 Neural Networks: Discriminator & Generator


Input: random (Gaussian) distribution
Output: an image

Random inputs are like latent representation,
Generator is like a decoder in VAE

Input: Either a fake image from the generator or a real image from the training set, and must guess whether the input image is fake or real.



GANs: Generative Adversarial Networks



The generator never actually sees any real images! But it gradually learns to produce convincing fake images guided by the gradients flowing back through the discriminator.


2. Training of the discriminator. A batch of real images is sampled from the training set and is completed with an equal number of fake images produced by the generator. The labels are set to 0 for fake images and 1 for real images, and the discriminator is trained on this labeled batch for one step, using the binary cross-entropy loss.

1. Training of the generator: It is used to produce another batch of fake images, and once again the discriminator is used to tell whether the images are fake or real. This time we do not add real images in the batch, and all the labels are set to 1 (real): in other words, we want the generator to produce images that the discriminator will (wrongly) believe to be real! Crucially, the weights of the discriminator are frozen during this step, so backpropagation only affects the weights of the generator.

(Latent) Diffusion Models

<https://arxiv.org/abs/2112.10752>

<https://arxiv.org/abs/2102.09672>




(Latent) Diffusion Models

<https://arxiv.org/abs/2112.10752>

<https://arxiv.org/abs/2102.09672>

- Diffusion models transform an image of random noise into a target image
- To gradually denoise the image we use a U-Net
- To guide the denoising process we leverage a text encoder (e.g. CLIP)



Text-to-Image Diffusion Models (but also Img-2-Img, Inpainting etc)

"stained glass of darth vader, backlight, centered composition, masterpiece, photorealistic, 8k"



base image



mask image



generated image

- E.g. DALL-E3, Stable Diffusion, Midjourney, Imagen

"cat wizard, gandalf, lord of the rings, detailed, fantasy, cute, adorable, Pixar, Disney, 8k"




initial image



generated image

Trying to fix the limitations...

Panel 3: The wolf asking Little Red Riding Hood about the contents of the basket and where she is going.



Trying to fix the limitations...

ControlNets



Figure 1: Controlling Stable Diffusion with learned conditions. ControlNet allows users to add conditions like Canny edges (top), human pose (bottom), *etc.*, to control the image generation of large pretrained diffusion models. The default results use the prompt “a high-quality, detailed, and professional image”. Users can optionally give prompts like the “chef in kitchen”.

Life After Diffusion Models

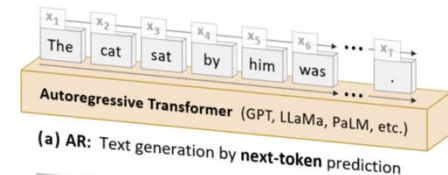
Autoregressive Image Generation

Probably:


- <https://arxiv.org/abs/2404.02905>
- <https://github.com/FoundationVision/VAR>




Three Different Autoregressive Generative Models



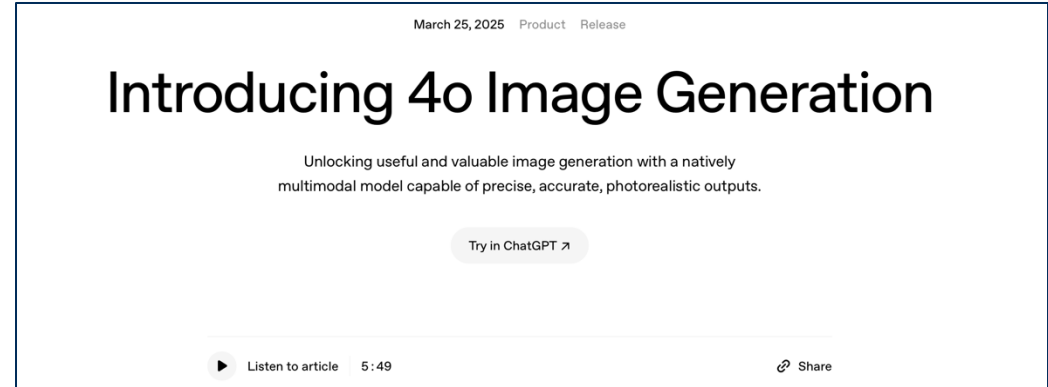
(a) AR: Text generation by next-token prediction



(b) AR: Image generation by next-image-token prediction



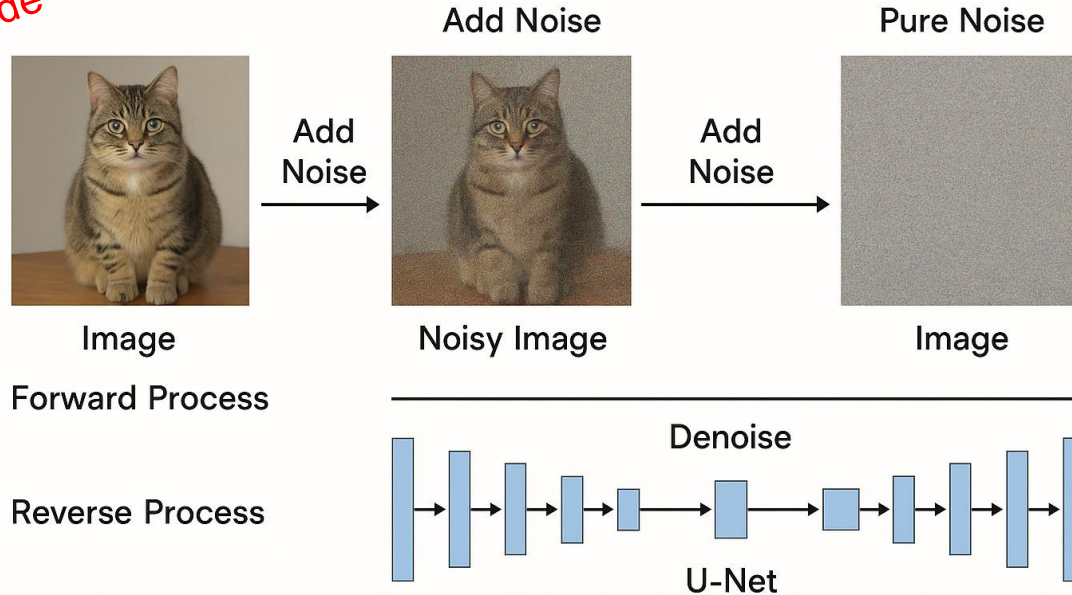
(c) VAR: Image generation by next-scale (or next-resolution) prediction



Diffusion Models explained by VisualAR (e.g GOT4o Image)

**CAREFUL!!! NOT Human
created Slide**

Diffusion Models



❖ Useful Reads:

- A. Géron, *Hands-on Machine Learning with Scikit-Learn, Keras, Tensorflow*, O'Reilly
- <https://www.kdnuggets.com/2021/03/beginners-guide-clip-model.html>
- https://huggingface.co/docs/transformers/main/en/tasks/zero_shot_object_detection
- https://huggingface.co/docs/transformers/main/en/tasks/object_detection
- <https://huggingface.co/docs/diffusers/index> (and following)
- https://huggingface.co/docs/transformers/model_doc/owlvit
- <https://towardsdatascience.com/stable-diffusion-using-hugging-face-501d8dbdd8#:~:text=Stable%20diffusion%20simply%20put%20is,image%20given%20a%20textual%20prompt.&text=As%20we%20can%20see%20from,image%20representative%20of%20the%20text.>
- <https://gregrobison.medium.com/tokens-not-noise-how-gpt-4os-approach-changes-everything-about-ai-art-99ab8ef5195d>
- And references within the presentation

Thank you.

